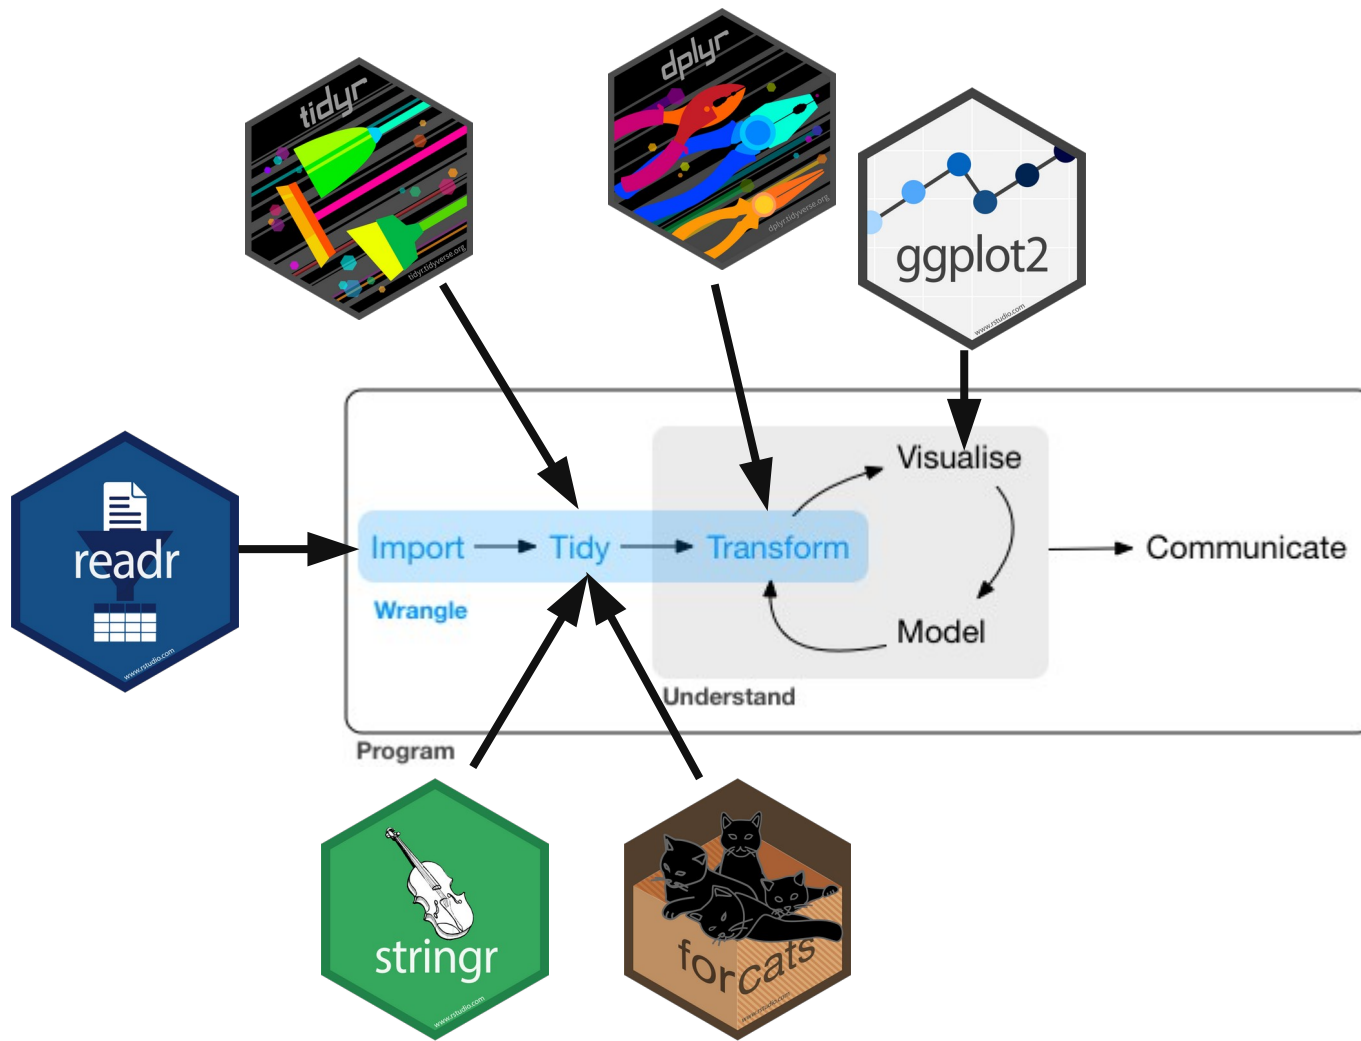


GG606

Factors & Dates





python: arrow, numpy, siuba, (pycats in pandas)

Factors

- Categorical variables: fixed, known set of possible values
- nb: Base R often converts characters to factors
- Read about `history` of `stringsAsFactors`

Factors

N

- Look-up table between Natural Numbers table of values
 - Example:
 - 1 = ctrl
 - 2 = trtA
 - 3 = trtB
- | |
|---------|
| 1 = Apr |
| 2 = Dec |
| 3 = Jan |

Nice Worked Example

- Two vectors
- `x1 ← c("Dec", "Apr", "Jan", "Mar")`
- `x2 ← c("Dec", "Apr", "Jam", "Mar")`
- `sort(x1)`
- Specify levels or use what's there?

```
sort(x1)
```

```
#> [1] "Apr" "Dec" "Jan" "Mar"
```

Nice Worked Example

- `month_levels ← c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")`
- `y1 ← factor(x1, levels = month_levels)`

```
y1
#> [1] Dec Apr Jan Mar
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
sort(y1)
#> [1] Jan Mar Apr Dec
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```

Nice Worked Example

- `month_levels ← c("Jan", "Feb", "Mar", "Apr", "May", "Jun", "Jul", "Aug", "Sep", "Oct", "Nov", "Dec")`
- `y1 ← factor(x1, levels = month_levels)`
- `y2 ← factor(x2, levels = month_levels)`

`y2`

```
#> [1] Dec Apr <NA> Mar
```

```
#> Levels: Jan Feb Mar Apr May Jun Jul Aug Sep Oct Nov Dec
```


Common (less safe?) Approach

- `factor(x1)`
- But only has some levels
- When would this matter or not matter?

```
factor(x1)
```

```
#> [1] Dec Apr Jan Mar
```

```
#> Levels: Apr Dec Jan Mar
```

Factors

- Look-up table between Natural Numbers
table of values

N

```
> as.numeric(y1)
[1] 12  4  1  3
> as.numeric(y2)
[1] 12  4 NA  3
```

General Social Survey

- Long-running US survey
- Good example of problems/challenges
- `forcats :: gss_cat`
- Not dissimilar to census data

gss_cat

- 9 columns
- Integers and factors
- `glimpse(gss_cat)`
- `str(gss_cat)`

```


> glimpse(gss_cat)
Rows: 21,483
Columns: 9
$ year      <int> 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2000, 2...
$ marital   <fct> Never married, Divorced, Widowed, Never married, Divorced, Married, Never marri...
$ age       <int> 26, 48, 67, 39, 25, 25, 36, 44, 44, 47, 53, 52, 52, 51, 52, 40, 77, 44, 40, 45,...
$ race      <fct> White, White, White, White, White, White, White, White, White, White, White, White, Wh...
$ rincome   <fct> $8000 to 9999, $8000 to 9999, Not applicable, Not applicable, Not applicable, $...
$ partyid   <fct> "Ind,near rep", "Not str republican", "Independent", "Ind,near rep", "Not str d...
$ relig     <fct> Protestant, Protestant, Protestant, Orthodox-christian, None, Protestant, Chris...
$ denom     <fct> Southern baptist, Baptist-dk which, No denomination, Not applicable, Not applic...
$ tvhours   <int> 12, NA, 2, 4, 1, NA, 3, NA, 0, 3, 2, NA, 1, NA, 1, 7, NA, 3, 3, NA, 1, 2, 2, 1,...
> str(gss_cat)
tibble [21,483 × 9] (S3: tbl_df/tbl/data.frame)
 $ year      : int [1:21483] 2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 2000 ...
 $ marital: Factor w/ 6 levels "No answer","Never married",..: 2 4 5 2 4 6 2 4 6 6 ...
 $ age       : int [1:21483] 26 48 67 39 25 25 36 44 44 47 ...
 $ race      : Factor w/ 4 levels "Other","Black",..: 3 3 3 3 3 3 3 3 3 3 3 ...
 $ rincome: Factor w/ 16 levels "No answer","Don't know",..: 8 8 16 16 16 5 4 9 4 4 ...
 $ partyid: Factor w/ 10 levels "No answer","Don't know",..: 6 5 7 6 9 10 5 8 9 4 ...
 $ relig    : Factor w/ 16 levels "No answer","Don't know",..: 15 15 15 6 12 15 5 15 15 15 ...
 $ denom    : Factor w/ 30 levels "No answer","Don't know",..: 25 23 3 30 30 25 30 15 4 25 ...
 $ tvhours  : int [1:21483] 12 NA 2 4 1 NA 3 NA 0 3 ...

```

Check factors

- `gss_cat %>%
 count(race)`
- `ggplot(gss_cat, aes(race)) +
 geom_bar()`
-

```
# A tibble: 3 x 2  
  race      n  
* <fct> <int>  
1 Other   1959  
2 Black   3129  
3 White  16395
```

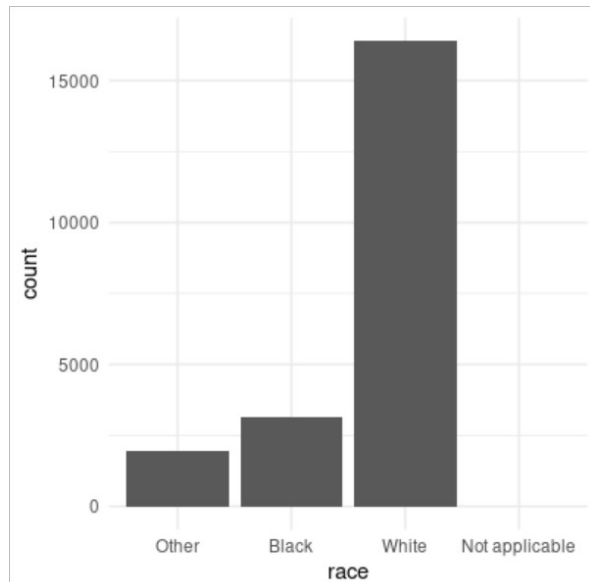



Check factors

- `gss_cat %>%
 count(race)`
- `ggplot(gss_cat, aes(race)) +
 geom_bar()`
- `ggplot(gss_cat, aes(race)) +
 geom_bar() +
 scale_x_discrete(drop = FALSE)`

A tibble: 3 x 2

	race	n
1	Other	1959
2	Black	3129
3	White	16395



Try this

- Explore the distribution of `rincome` (reported income). What makes the default bar chart hard to understand? How could you improve the plot?
- What is the most common `relig` in this survey? What's the most common `partyid`?

Factor order

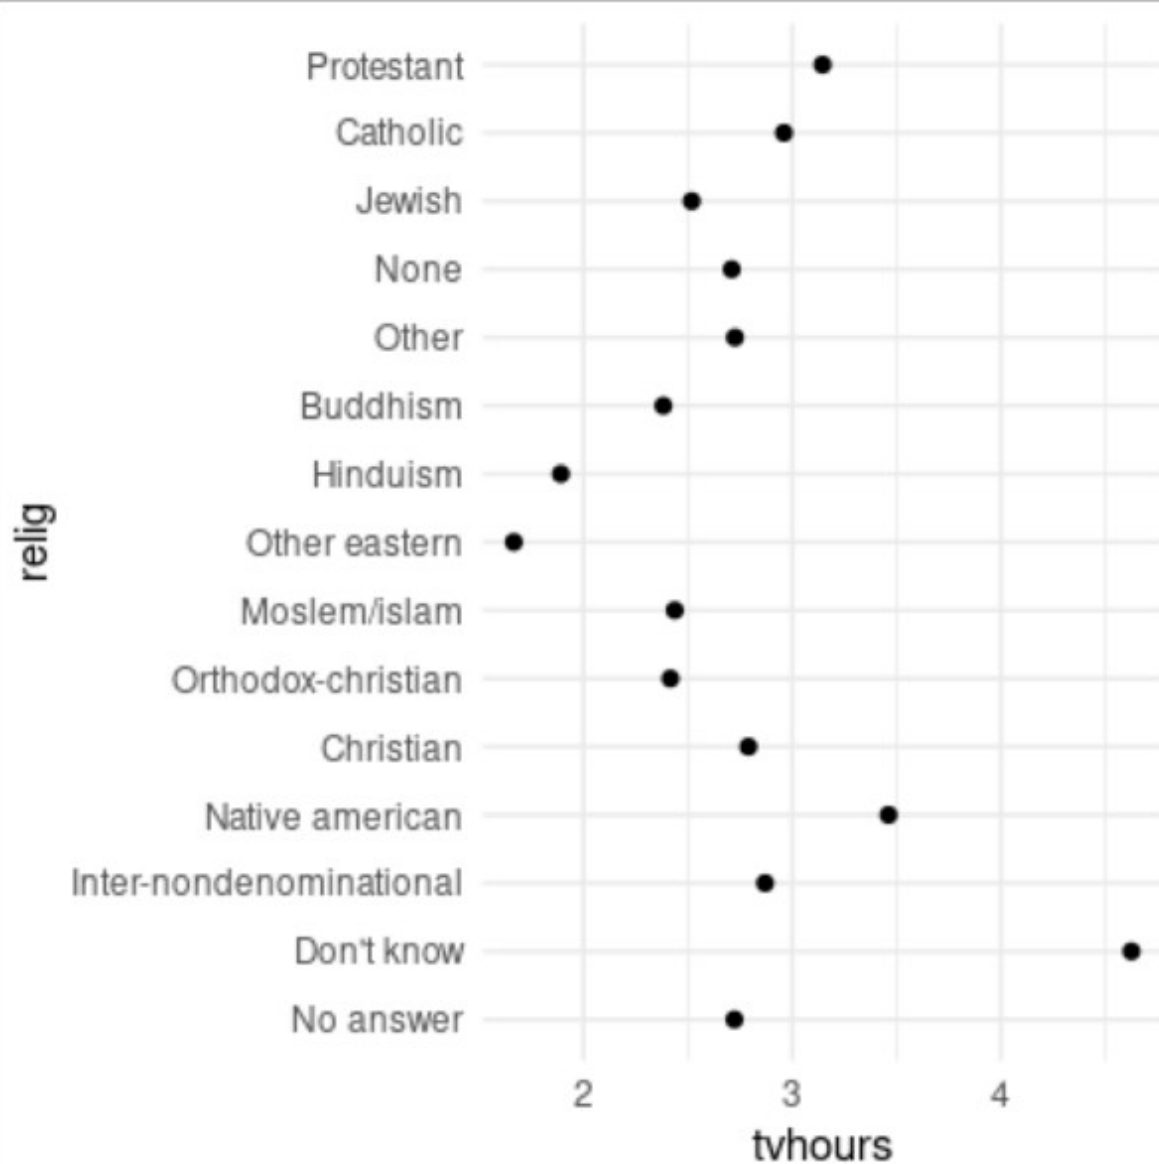
- Average hours watching TV per religion
- Hint: use `group_by` & `summarise`

Factor order

- ```
relig_summary ← gss_cat %>%
 group_by(relig) %>%
 summarise(
 age = mean(age, na.rm = TRUE),
 tvhours = mean(tvhours, na.rm = TRUE),
 n = n()
)
```
- ```
ggplot(relig_summary, aes(tvhours, relig)) +  
  geom_point()
```

Fact

- `relig_summary ← gss_group_by(relig) %>% summarise(
 age = mean(age, n
 tvhours = mean(tv
 n = n()
)`
- `ggplot(relig_summary,
 geom_point()`



Factor reorder


- `fct_reorder(f, x)`
- Defaults to median, can specify functions
 - Reorder `f` according to `x`

Factor reorder

- `fct_reorder(f, x)`
- Defaults to median, can specify functions
- `ggplot(relig_summary, aes(tvhours, fct_reorder(relig, tvhours))) +
 geom_point()`
- Better to do it with `mutate` or outside of `ggplot`

Factor

- `fct_reorder(f,`
- `ggplot(relig_s`
`fct_reorder(re`
`geom_point()`



Factor reorder

- Income vs age?
- ```
rincome_summary <- gss_cat %>%
 group_by(rincome) %>%
 summarise(
 age = mean(age, na.rm = TRUE),
 tvhours = mean(tvhours, na.rm = TRUE),
 n = n()
)
```
- ```
ggplot(rincome_summary, aes(age, fct_reorder(rincome, age))) +  
  geom_point()
```
- Does it make any sense?

Other reorder

- `fct_reorder2`: orders by y with largest x (2D)
- `fct_infreq`: orders by frequency
- ```
gss_cat %>%
 mutate(marital = marital %>%
 fct_infreq() %>% fct_rev()) %>%
 ggplot(aes(marital)) +
 geom_bar()
```



# Modifying levels

- `fct_recode`: change (recode) values of levels
- `gss_cat %>% count(partyid)`

```
A tibble: 10 x 2
 partyid n
 <fct> <int>
1 No answer 154
2 Don't know 1
3 Other party 393
4 Strong republican 2314
5 Not str republican 3032
6 Ind,near rep 1791
7 Independent 4119
8 Ind,near dem 2499
9 Not str democrat 3690
10 Strong democrat 3490
```

# Modifying levels

- `gss_cat %>%`

```
 mutate(partyid = fct_recode(partyid,
 "Republican, strong" = "Strong republican",
 "Republican, weak" = "Not str republican",
 "Independent, near rep" = "Ind,near rep",
 "Independent, near dem" = "Ind,near dem",
 "Democrat, weak" = "Not str democrat",
 "Democrat, strong" = "Strong democrat"
)) %>%
count(partyid)
```

# A tibble: 10 x 2

|    | partyid<br>* <fct> |
|----|--------------------|
| 1  | No answer          |
| 2  | Don't know         |
| 3  | Other party        |
| 4  | Strong republican  |
| 5  | Not str republican |
| 6  | Ind,near rep       |
| 7  | Independent        |
| 8  | Ind,near dem       |
| 9  | Not str democrat   |
| 10 | Strong democrat    |

# A tibble: 10 x 2

| n    | partyid<br><int> <fct> |
|------|------------------------|
| 154  | No answer              |
| 1    | Don't know             |
| 393  | Other party            |
| 2314 | Republican, strong     |
| 3032 | Republican, weak       |
| 1791 | Independent, near rep  |
| 4119 | Independent            |
| 2499 | Independent, near dem  |
| 3690 | Democrat, weak         |
| 3490 | Democrat, strong       |

| n    | <int> |
|------|-------|
| 154  |       |
| 1    |       |
| 393  |       |
| 2314 |       |
| 3032 |       |
| 1791 |       |
| 4119 |       |
| 2499 |       |
| 3690 |       |
| 3490 |       |

# Collapsing levels

- ```
gss_cat %>%  
  mutate(partyid = fct_collapse(partyid,  
    other = c("No answer", "Don't know", "Other party"),  
    rep = c("Strong republican", "Not str republican"),  
    ind = c("Ind,near rep", "Independent", "Ind,near dem"),  
    dem = c("Not str democrat", "Strong democrat")  
  )) %>%  
  count(partyid)
```
- See also `?fct_lump`





When does/did time begin?
When is $t=0$?

Disney · PIXAR
FORKY ASKS A QUESTION:

WHAT IS
TIME?

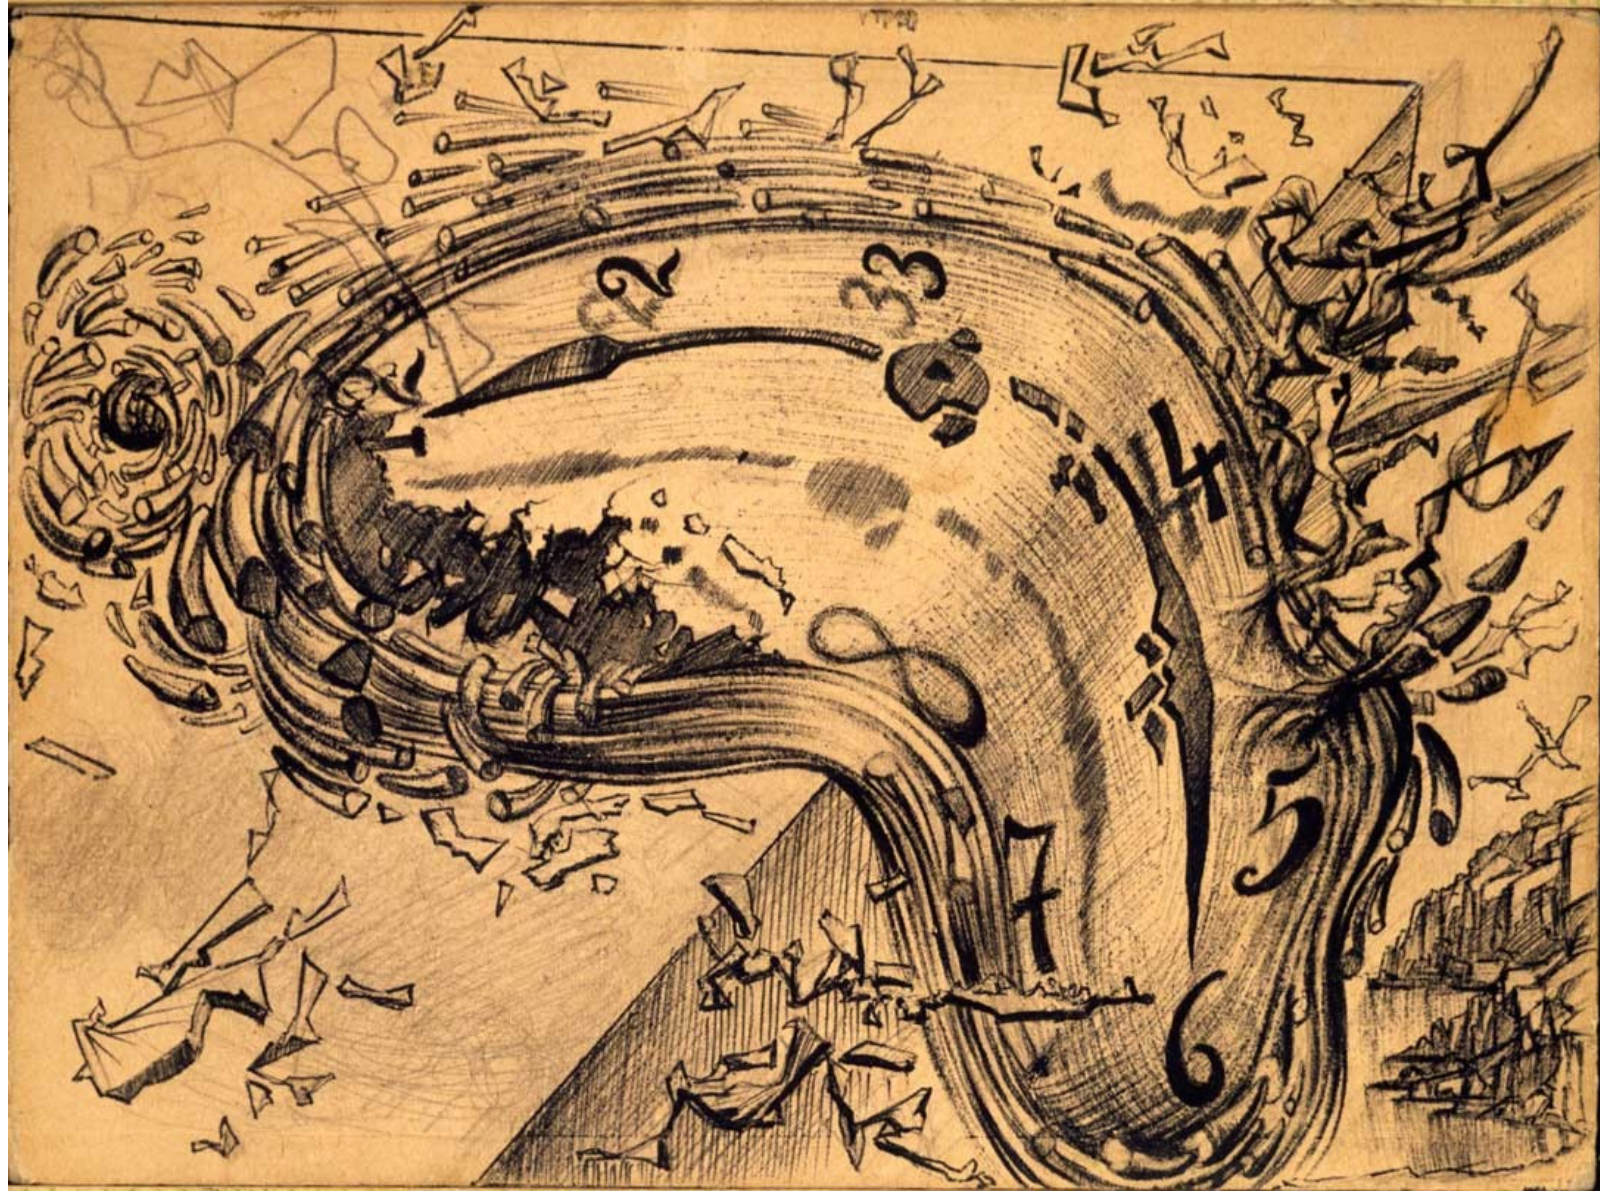




When does/did time begin?

When is $t=0$?

Open spreadsheet, type 0, change
format to YYYY-MM-DD hh:mm:ss



lubridate

- Friendly with `tidyverse`
- `ts`, `xts` & `zoo` for time series
 - time + matrix
 - Common for predictions, some models

How is time counted

- Excel: days since 1900-01-01 or 1904-01-01
- POSIX: seconds since 1970-01-01 UTC
 - POSIXct: (calendar/continuous time) seconds
 - POSIXlt: (local time) list with elements
- What gets exported to csv?

Types

- date
- time
- date-time <dtm>: POSIXct
- Who types in date-time format in spreadsheets?

Construct Dates

- `today()` & `now()`
- Helper functions for numbers & strings:

- `ymd(20200708)`
`ymd("2020-07-08")`
`ymd("2020-Jul-08")`
- `mdy("July 8, 2020")`
- `dmy("08-July-2020")`
`dmy("08-07-2020")`

```
> ymd(20200708)
[1] "2020-07-08"
> ymd("2020-07-08")
[1] "2020-07-08"
> ymd("2020-Jul-08")
[1] "2020-07-08"
> mdy("July 8, 2020")
[1] "2020-07-08"
> dmy("08-July-2020")
[1] "2020-07-08"
> dmy("08-07-2020")
[1] "2020-07-08"
```

Construct Dates & Times

- `ymd_hms("2020-07-08 20:11:59")`
- `mdy_hm("07/08/2017 08:01 pm")`
- Converts to UTC internally

```
> ymd_hms("2020-07-08 20:11:59")
```

```
[1] "2020-07-08 20:11:59 UTC"
```

```
> mdy_hm("07/08/2017 08:11 pm")
```

```
[1] "2017-07-08 20:11:00 UTC"
```

Construct Dates & Times

- Converts to UTC internally
- `ymd_hms("2020-07-08 20:11:59") %>% as.POSIXct()`
- `ymd_hms("2020-07-08 20:11:59", tz="EDT") %>% as.POSIXct()`

```
> ymd_hms("2020-07-08 20:11:59") %>% as.POSIXct()  
[1] "2020-07-08 20:11:59 UTC"  
> ymd_hms("2020-07-08 20:11:59", tz="EDT") %>% as.POSIXct()  
[1] "2020-07-09 00:11:59 EDT"
```

Time Zones & locales

- `Sys.getlocale("LC_TIME")`

```
> Sys.getlocale("LC_TIME")  
[1] "en_CA.UTF-8"
```

`tz` a character string that specifies which time zone to parse the date with. The string must be a time zone that is recognized by the user's OS.

Another problem is that name needs to reflect not only to the current behaviour, but also the *complete history*. For example, there are time zones for both “America/New_York” and “America/Detroit”. These cities both currently use Eastern Standard Time but in 1969-1972 *Michigan (the state in which Detroit is located), did not follow DST*, so it needs a different name.

Time Zones

- Complicated
- IANA time zones:
<continent>/<city>
 - "Canada/Eastern"
 - "America/Toronto"
 - "America/Nipigon"
 - "America/Thunder_Bay"
- `Sys.timezone()`
- `OlsonNames()`

#	Zone	NAME	STDOFF	RULES	FORMAT	[UNTIL]
	Zone	America/Toronto	-5:17:32	-	LMT	1895
			-5:00	Canada	E%sT	1919
			-5:00	Toronto	E%sT	1942 Feb 9 2:00s
			-5:00	Canada	E%sT	1946
			-5:00	Toronto	E%sT	1974
			-5:00	Canada	E%sT	
	Zone	America/Thunder_Bay	-5:57:00	-	LMT	1895
			-6:00	-	CST	1910
			-5:00	-	EST	1942
			-5:00	Canada	E%sT	1970
			-5:00	Toronto	E%sT	1973
			-5:00	-	EST	1974
			-5:00	Canada	E%sT	
	Zone	America/Nipigon	-5:53:04	-	LMT	1895
			-5:00	Canada	E%sT	1940 Sep 29
			-5:00	1:00	EDT	1942 Feb 9 2:00s
			-5:00	Canada	E%sT	
	Zone	America/Rainy_River	-6:18:16	-	LMT	1895
			-6:00	Canada	C%sT	1940 Sep 29
			-6:00	1:00	CDT	1942 Feb 9 2:00s
			-6:00	Canada	C%sT	
	Zone	America/Atikokan	-6:06:28	-	LMT	1895
			-6:00	Canada	C%sT	1940 Sep 29
			-6:00	1:00	CDT	1942 Feb 9 2:00s
			-6:00	Canada	C%sT	1945 Sep 30 2:00
			-5:00	-	EST	

Time Zones Control Printing

- `(x1 ← ymd_hms("2015-06-01 12:00:00", tz = "America/New_York"))`
- `(x2 ← ymd_hms("2015-06-01 18:00:00", tz = "Europe/Copenhagen"))`
- `(x3 ← ymd_hms("2015-06-02 04:00:00", tz = "Pacific/Auckland"))`
- `x1 - x2`

```
> (x1 ← ymd_hms("2015-06-01 12:00:00", tz = "America/New_York"))  
[1] "2015-06-01 12:00:00 EDT"  
> (x2 ← ymd_hms("2015-06-01 18:00:00", tz = "Europe/Copenhagen"))  
[1] "2015-06-01 18:00:00 CEST"  
> (x3 ← ymd_hms("2015-06-02 04:00:00", tz = "Pacific/Auckland"))  
[1] "2015-06-02 04:00:00 NZST"  
> x1 - x2  
Time difference of 0 secs
```
- `x1 - x3`

```
> x1 - x3  
Time difference of 0 secs
```